CSUF

COLLEGE OF

Engineering and Computer Science

Predicting GDP Growth Using Economic Indicators

Department of Computer Science

Team: Zeid Aldaas, Salmaan Rauf, Abdurrahman Balsus

Team Number: 3

Faculty Advisor: Dr. Yu Bai

TABLE OF CONTENTS

I. Abstract 1

II. Introduction 1

III. Methodology 4

IV. Results 6

V. Discussion 8

VI. Conclusion 9

VII. References 11

I. Abstract

This project aims to predict the GDP growth rate of the United States using neural networks, leveraging key economic indicators as predictors. By modeling GDP growth, policymakers and analysts can gain insights into how inflation, unemployment, interest rates, government spending, imports, foreign direct investment (FDI), domestic credit to the private sector, and population growth affect economic performance. Data were obtained from the World Bank's World Development Indicators database and subjected to preprocessing steps that addressed missing values and scale differences. A multi-layer perceptron was trained and evaluated against standard regression metrics. While initial linear regression models performed poorly, the neural network approach achieved a coefficient of determination (R^2) of approximately 0.75 and a mean squared error (MSE) of 0.49. Exploratory data analysis revealed that private domestic credit was the most influential predictor, while population growth and imports were less impactful. The final model underscores the potential of neural networks in macroeconomic forecasting, yet the complexities of multicollinearity, missing indicators, and unpredictable global events present limitations. Future work may expand this approach to other countries, explore more sophisticated deep learning architectures, and incorporate additional explanatory variables to enhance both accuracy and interpretability.

II. Introduction

Economic growth, often quantified as the annual percentage change in Gross Domestic Product (GDP), is a cornerstone metric for assessing a country's overall economic health and stability. Accurate GDP growth forecasting can inform policy decisions, guide investment strategies, and support macroeconomic planning. Traditional forecasting methods, such as linear regression, have provided baseline insights but frequently fail to capture the complex, nonlinear relationships that characterize real-world economies [1]. Consequently, the use of machine learning (ML) techniques, particularly neural networks, has gained traction in economic research and policy analysis due to their ability to model intricate interactions among variables.

Background and Significance:

This project addresses the challenge of modeling and predicting GDP growth rates using a combination of well-established economic indicators. Key indicators considered include inflation rate, unemployment rate, interest rate, government spending, imports, FDI, private domestic credit, and population growth. These factors are widely cited in the literature as influential drivers of economic outcomes. For instance, expansions in credit markets have been linked to economic growth, while high unemployment often signals underlying structural issues that constrain development [2]–[4]. Inflation and interest rates are critical for monetary policy considerations, and government spending patterns can stimulate or dampen growth. Imports and foreign direct investment have extensive literature exploring their roles in both short-term fluctuations and long-term development trajectories [5]–[7]. Population growth, while often viewed through a demographic lens, can also influence labor force dynamics and productivity.

Literature Review and Gaps:

Prior studies on GDP forecasting have used various methodologies ranging from purely econometric models (e.g., Vector Autoregression models) to hybrid techniques that incorporate time-series decomposition and machine learning. Traditional linear regression models offer simplicity and interpretability but may fail to capture the nonlinearities and interactions present in macroeconomic data. More advanced ML approaches—such as Random Forests, Gradient Boosted Trees, and Artificial Neural Networks—have shown promise in improving predictive accuracy [8].

Neural networks, in particular, are advantageous due to their flexibility and ability to approximate complex functions. Studies have shown that neural networks can outperform traditional econometric methods under certain conditions, especially when sufficient high-quality data is available and the complexity of the economic system is high [9]. However, challenges persist: macroeconomic data often suffer from missing values, inconsistent reporting standards, and multicollinearity among indicators. Moreover, unforeseen external shocks—global crises, pandemics, geopolitical changes—introduce non-stationarities and limit the predictive power of any model trained on historical data.

This study attempts to address two primary gaps:

1. **Integration of Multiple Indicators:** Instead of relying on a few macroeconomic variables, this project incorporates a diverse set of predictors, thereby offering a more holistic perspective on GDP growth drivers.

2. **Application of Neural Networks for U.S. GDP Forecasting:** While neural networks have been tested for macroeconomic forecasting, their application to U.S. data with a comprehensive set of indicators and a clear focus on model interpretability remains an area ripe for exploration.

Objectives:

The primary objective of this project is to develop a predictive model for U.S. GDP growth using neural networks and to identify which economic indicators are most influential. Specific goals include:

- Collecting and preprocessing data from the World Bank's World Development Indicators to ensure consistency and quality.

- Comparing a neural network-based model's performance against baseline methods (initially linear regression) to assess improvement in predictive accuracy.

- Evaluating feature importance and discussing the interpretability of the results.

- Identifying model limitations and suggesting directions for future research, including scaling the approach to other countries and incorporating more advanced ML techniques.

Scope and Limitations:

This project focuses on the United States during a historical period covered by the available World Bank data. Limitations include potential data quality issues, missing indicators, and the inability to account for sudden, exogenous shocks such as global financial crises or pandemics. While the model may generalize well to similar advanced economies, the uniqueness of the U.S. economic environment might limit its direct applicability elsewhere. Nonetheless, the methodologies and insights gleaned can serve as a valuable template for future studies in other contexts.

III. Methodology

The methodology for this project involved data acquisition, preprocessing, exploratory data analysis (EDA), model development, and evaluation. A schematic of the workflow is presented in Figure 1.

Data Source:

Data were obtained from the World Bank's World Development Indicators (WDI) database [10]. The selected variables included:

- **Target Variable:** GDP growth (annual %)
- **Predictors:** Unemployment rate, Imports (% of GDP), FDI (net outflows, % of GDP), Domestic credit to the private sector (% of GDP), Population growth (annual %), Inflation (annual %), Real interest rate (%), and Government consumption expenditure (% of GDP).

Data Preprocessing:

1. **Data Cleaning:**

    - Dropped unnecessary columns (e.g., country names and codes) not required for modeling.

    - Converted wide format data into a long format and then pivoted back to ensure each year became a unique entry and each economic indicator formed a column.

    - Handled missing values by removing rows containing NaNs to maintain data integrity.

2. **Feature Engineering and Selection:**

    Given the scope and the established literature, no additional feature engineering was performed beyond selecting the key economic indicators. Predictors were chosen based on theoretical importance and data availability.

3. **Scaling and Encoding:**

    A **StandardScaler** transformation was applied to all predictor variables to ensure comparability and to prevent features with large numeric ranges from dominating the model. No categorical encoding was needed since all chosen predictors were numeric.

Modeling Approach:

Initially, a multilinear regression model was tested, but results indicated poor predictive performance. To capture nonlinear relationships, the approach moved to a neural network architecture (Multi-Layer Perceptron, MLP) implemented using Python's scikit-learn. Two different MLP regressors were combined into a stacking ensemble:

- **MLP1:** Hidden layers: (512, 256), max_iter=3000

- **MLP2:** Hidden layers: (256, 128, 64), max_iter=3000

- **Stacking Regressor:** Final estimator MLP with hidden layers (128, 64), max_iter=3000

This stacked approach aimed to leverage the diversity of multiple neural network configurations, potentially improving robustness and generalization.

Training and Validation:

The data were split into training and test sets (80% training, 20% test) using a random state for reproducibility. The model was trained to minimize MSE, and hyperparameters were chosen based on default settings and iteration limits. No extensive hyperparameter tuning (e.g., grid search) was performed due to project constraints.

Evaluation Metrics:

Performance was measured using:

- **Coefficient of Determination (R^2):** Indicates the proportion of variance explained by the model.
- **Mean Squared Error (MSE):** Measures average squared difference between predicted and actual values.

Interpretability Tools:

Permutation feature importance was applied to assess the relative importance of each predictor, facilitating insights into which indicators most influenced the model's predictions.

IV. Results

The final stacked neural network model achieved an R^2 of 0.75 and an MSE of 0.49 on the test set. These metrics reflect a substantial improvement over the initial linear regression attempt, which struggled to provide meaningful forecasts.

Quantitative Outcomes:

- **R^2 Score:** 0.75
- **MSE:** 0.49

The R^2 score indicates that the model captured approximately three-quarters of the variance in GDP growth, while the MSE suggests moderately accurate predictions, given the complexity of economic data.

Visual Representations:

Figure 2: A correlation heatmap displayed the relationships among the target (GDP growth) and all predictors. Strong correlations were observed between private domestic credit and GDP growth, while weaker or more complex relationships were seen for population growth and imports.

Figure 3: Actual vs. Predicted GDP growth distributions were plotted using kernel density estimation (KDE). The distributions were reasonably aligned, indicating that the model's predictions closely approximate the underlying data's shape.

Figure 4: Permutation feature importance indicated that private domestic credit (% of GDP) emerged as the most impactful predictor. Inflation, unemployment, and interest rates followed. Conversely, population growth and imports had comparatively lower importance scores, suggesting limited direct influence on short-term GDP growth predictions.

Figure 5: Residuals plotted over time showed no clear temporal trends, although there were periods with slightly larger deviations. The residual analysis suggested that while the model performed well overall, certain historical contexts or outlier years might remain challenging to predict accurately.

V. Discussion

The results reinforce the hypothesis that a neural network-based model can outperform a simple linear regression approach for predicting GDP growth. The relatively high R^2 score (0.75) demonstrates that the chosen predictors and model architecture captured a significant portion of the variance in GDP growth rates. However, economic forecasting remains inherently challenging due to factors like multicollinearity, missing indicators, and the unpredictable nature of geopolitical and global health events.

Interpretation of Results:

- **Private Domestic Credit:** Its high feature importance suggests that credit availability and financial development play a pivotal role in driving economic growth. This finding aligns with economic theory, where well-functioning credit markets facilitate investment and consumer spending.
- **Limited Importance of Population Growth and Imports:** While intuitively important, these variables may be overshadowed by more direct financial and policy-related drivers or may require additional context (such as productivity measures) to influence GDP growth predictions.

Limitations:

1. **Data Quality and Coverage:** Missing values and limited availability of certain indicators constrained the dataset.
2. **Multicollinearity:** Economic indicators often move together, complicating attribution of importance and interpretability.
3. **External Shocks:** The model did not explicitly account for sudden events (e.g., COVID-19 pandemic, global financial crises) that dramatically alter economic trajectories.

4. **Temporal Stability:** Relationships identified in one historical period may not hold in the future, limiting the model's forecasting horizon.

Future Research Directions:

- **Geographical Expansion:** Applying this methodology to different countries or regional blocs may provide more generalizable insights.
- **Model Complexity:** Incorporating more sophisticated deep learning architectures (e.g., Long Short-Term Memory networks for time-series data) might improve predictive performance.
- **Additional Indicators:** Introducing sector-specific data, productivity measures, or geopolitical stability indices could enhance model accuracy and explainability.
- **Robustness Checks:** Employing cross-validation, Bayesian optimization for hyperparameter tuning, or ensembling with other models could further refine results.

VI. Conclusion

This project demonstrated the feasibility and utility of using neural networks to predict U.S. GDP growth from a diverse set of economic indicators. By leveraging data from the World Bank and focusing on key variables (inflation, unemployment, interest rates, government spending, imports, FDI, private domestic credit, and population growth), the model achieved an R^2 of 0.75, signifying robust predictive capabilities compared to a simple linear approach.

The analysis highlighted private domestic credit as a crucial driver of GDP growth, while population growth and imports exerted relatively weaker influences. Despite these advances, challenges in data quality, interpretability, and susceptibility to external shocks remain. The project's approach and findings contribute to the literature on machine learning applications in macroeconomics, offering a template for more extensive future work that can incorporate additional countries, richer datasets, and cutting-edge modeling techniques.

VII. References

[1] J. Stock and M. Watson, "Forecasting Output and Inflation: The Role of Asset Prices," *J. Econ. Lit.*, vol. 41, no. 3, pp. 788–829, 2003.

[2] R. Levine, "Financial Development and Economic Growth: Views and Agenda," *J. Econ. Lit.*, vol. 35, no. 2, pp. 688–726, 1997.

[3] G. Cecchetti and E. Kharroubi, "Reassessing the Impact of Finance on Growth," *BIS Working Papers*, no. 381, 2012.

[4] A. Afonso and J. González Alegre, "Economic Growth and Budgetary Components: A Panel Assessment for the EU," *Empirica*, vol. 38, no. 3, pp. 505–518, 2011.

[5] N. Borensztein, J. De Gregorio, and J.-W. Lee, "How Does Foreign Direct Investment Affect Economic Growth?" *J. Int. Econ.*, vol. 45, no. 1, pp. 115–135, 1998.

[6] P. F. Druck and P. Tambakis, "Global Economic Conditions and Central Bank Policy Responses," *Int. Monet. Fund*, 2015.

[7] P. Romer, "Endogenous Technological Change," *J. Polit. Econ.*, vol. 98, no. 5, pp. S71–S102, 1990.

[8] R. L. Trapletti, S. Hornik, and K. Hornik, "Forecasting Economic Time Series with Neural Networks: The Case of Austrian Tourism Data," *J. Forecast.*, vol. 19, no. 8, pp. 605–618, 2000.

[9] G. Zhang, B. Eddy Patuwo, and M. Y. Hu, "Forecasting with Artificial Neural Networks: The State of the Art," *Int. J. Forecast.*, vol. 14, no. 1, pp. 35–62, 1998.

[10] The World Bank Group, "World Development Indicators," [Online]. Available: https://data.worldbank.org/indicator